

Bhanuprasad Kokkula

New York, NY | (848) 391-1376 | bhanu.kokkula457@gmail.com
linkedin.com/in/bhanu-kokkula | github.com/Bhanukokkula | bhanukokkula.github.io

EDUCATION & CERTIFICATIONS

Rutgers, The State University of New Jersey — M.S. in Data Science Aug 2024 – May 2026
University of Pennsylvania (Wharton School) — Professional Certificate: AI for Business Jan 2026 – Jun 2026
Certifications: Bloomberg Market Concepts · Microsoft Certified: Azure Data Scientist Associate
Selected Coursework: Machine Learning Summer Program – IIT Kharagpur

EXPERIENCE

Rutgers, The State University of New Jersey — Graduate GenAI Engineer Oct 2024 – May 2026

- Built RAG pipelines and fine-tuned LLaMA models with LoRA/PEFT for an enterprise LLM assistant, running R&D across FAISS, Pinecone, and ChromaDB vector stores and integrating OpenAI and Anthropic (Claude) APIs for ~50% higher task precision over the base model.
- Built day to day with AI coding assistants (Claude, GitHub Copilot) and an automated evaluation harness benchmarking 15+ model and retrieval configurations on accuracy, grounding, and latency, cutting evaluation cycles from 2 weeks to 3 hours.
- Ran controlled prompt and retrieval experiments across 5K+ test cases and shipped the validated configuration as the team default, lifting answer accuracy 31% over baseline.

Hexaware Technologies — Data Scientist, Financial Services Jun 2023 – Aug 2024

- Built and integrated REST APIs to serve model outputs and connect data services across client systems, backing production fraud and credit-risk models over 5M+ records on AWS SageMaker.
- Owned feature pipelines over 500K+ daily transactions with automated validation, drift monitoring, and retraining; features drove model decisions that cut fraud losses 8%.
- Partnered directly with client and business stakeholders to turn requirements into deployed solutions, automating 14 manual workflows and cutting turnaround from 3 days to 4 hours.

Chinmaya Technologies — Data Science Intern Jan 2023 – May 2023

- Built distributed data pipelines in PySpark on Databricks processing 1M+ transaction and clickstream records, engineering 40+ behavioral features for downstream models.
- Delivered SKU-level demand forecasts that cut stockouts 12%, sharing results in recurring reports that guided supply-chain planning decisions.

Echortech — Data Engineering Intern Sep 2022 – Dec 2022

- Built and deployed ETL pipelines with Airflow moving millions of records daily from 6 source systems into BigQuery, with automated data-quality checks.

SELECTED PROJECTS

MCP Codebase Operator — Safe Agent Tooling Server Python, MCP, Claude Code, GitHub Actions

- Built an MCP server in Python that exposes a codebase to AI agents through guarded tools (test runner, database migrations, log access), enforcing read/write separation and confirmation gates on destructive actions so an agent operates safely instead of running unchecked.
- Connected it to Claude Code over MCP with GitHub webhooks and Actions, logging every tool call for a full audit trail and blocking any change that failed the test suite so no untested code merged.

MARS — Multi-Agent Research & Ingestion Platform Python, LangGraph, ChromaDB, FastAPI, Docker

- Built a multi-agent LLM system with LangGraph and a ChromaDB RAG layer synthesizing SEC filings, macro data, and news across 50K+ documents, coordinating agents to extract, score, and reconcile cross-source signals.
- Containerized the ingestion pipelines and FastAPI services with Docker for reproducible deployment, backed by 123 passing tests, cutting research turnaround from 6 hours to 15 minutes.

COVENANT — Document Extraction & Validation Pipeline Python, Pydantic, SQLite, Streamlit

- Built an LLM pipeline extracting structured data from documents (Pydantic + SQLite) with a Streamlit interface, shipping a working proof of concept and iterating with human-in-the-loop QA.
- Wired GitHub Actions to run the full suite of 97 tests on every commit before shipping, reaching 93% Q&A accuracy on the governed dataset.

RESEARCH & PUBLICATION

Rutgers, The State University of New Jersey — Research Co-author, Computational Linguistics (ongoing) Jun 2025

- Built a text-mining pipeline processing 194K+ abstracts and 69,580 full-text papers with semantic claim-matching, applying journal-clustered logistic regressions and pre-registered behavioral trials (N = 1,105) to quantify statistical claim inflation.

SKILLS

Languages: Python (scripting, automation, API integration), SQL, Bash, TypeScript/JavaScript, Node.js
AI Dev Tools & Agents: Claude Code, GitHub Copilot, MCP (servers/clients), LangGraph, LangChain, Prompt Engineering, Agentic Workflows
LLMs & RAG: LLM Fine-Tuning (LoRA/PEFT), RAG, Embeddings, Hugging Face; OpenAI, Anthropic (Claude), LLaMA; Vector Stores (ChromaDB, FAISS, Pinecone)
Integration & Connectivity: REST APIs, Webhooks, MCP, Git/GitHub Workflows, CI/CD
DevOps & Cloud: Docker, GitHub Actions, Linux, MLflow; AWS, GCP, Azure